

Nested sampling

An efficient and robust Bayesian inference tool for 21cm cosmology

Will Handley

<wh260@cam.ac.uk>

Royal Society University Research Fellow
Astrophysics Group, Cavendish Laboratory, University of Cambridge
Kavli Institute for Cosmology, Cambridge
Gonville & Caius College
willhandley.co.uk

21st October 2020



UNIVERSITY OF
CAMBRIDGE



Bayesian inference in global 21cm cosmology

- ▶ Data analysis problem

$$\begin{aligned}\text{Data} &= \text{Signal} + \text{Foreground} + \text{Noise} \\ T_{\text{data}}(\nu) &= T_{21\text{cm}}(\nu; \theta_{21\text{cm}}) + T_{\text{fg}}(\nu; \theta_{\text{fg}}) + T_{\text{noise}}(\nu)\end{aligned}$$

- ▶ We can only statistically describe T_{noise}
 - ▶ e.g. as an (un)correlated Gaussian random variable $P(T_{\text{noise}}) = \frac{1}{\sqrt{2\pi}\sigma} e^{-T_{\text{noise}}^2/2\sigma^2}$
- ▶ Allows us to form a *likelihood* for the unknown parameters $\theta = (\theta_{21\text{cm}}, \theta_{\text{fg}})$

$$\begin{aligned}\mathcal{L}(\theta) &= P(\text{Data}|\theta) \propto e^{-\frac{1}{2}\chi(\theta)^2} \\ \chi^2(\theta) &= \sum_{\nu} \frac{1}{\sigma_{\nu}^2} [T_{\text{data}}(\nu) - T_{21\text{cm}}(\nu; \theta_{21\text{cm}}) - T_{\text{fg}}(\nu; \theta_{\text{fg}})]^2\end{aligned}$$

- ▶ This misses the initial step of data compression (reduction, flagging, integration)
- ▶ In practice also include θ_{noise} as parameters as well.

The three pillars of Bayesian inference

- ▶ Frequentist approaches maximise the likelihood $\mathcal{L}(\theta) = P(D|\theta)$
 - ▶ anything using a “least squares” optimisation (e.g. maxsmooth, EDGES calibration)
- ▶ Bayes theorem allows us to answer science questions using probabilities:

Parameter estimation Given a model, which range of parameters best describe the data?

$$P(\theta|D, M) = \frac{P(D|\theta, M)P(\theta|M)}{P(D|M)}, \quad \text{Posterior} = \frac{\text{Likelihood} \times \text{Prior}}{\text{Evidence}}$$

Model comparison Which models do the data prefer?

$$P(M|D) = \frac{P(D|M)P(M)}{P(D)}, \quad \text{Model Probability} = \frac{\text{Evidence} \times \text{Model Prior}}{\text{Normalisation}}$$


Tension quantification Are different (sub)sets of data consistent with one another?

Model comparison in 21cm cosmology

- ▶ Model comparison: Given this data, what odds would a bookmaker put on a model?
 - ▶ In general in 21cm cosmology model takes the form $T_{\text{data}} = T_{21\text{cm}} + T_{\text{fg}} + T_{\text{noise}}$
 - ▶ Model comparison can be applied to any combination of these pieces
1. Does the data contain a signal?
(Does $T_{\text{data}} = T_{21\text{cm}} + T_{\text{fg}} + T_{\text{noise}}$ have a higher evidence than $T_{\text{data}} = T_{\text{fg}} + T_{\text{noise}}$)
 2. Which foreground model is better? (polynomial vs forward sky model)
 3. Which noise model is better? (correlated gaussian, uncorrelated gaussian, lorentzian)
 4. How many components should a foreground model have?
 5. How many components should a calibration model have?


Numerical tools

- ▶ The key concept in numerical Bayesian inference is *sampling* a distribution
- ▶ i.e. drawing parameter vectors θ in proportion to the probability density $P(\theta)$.
- ▶ Compression of high-dimensional function.
- ▶ Posterior $P(\theta|D)$ encodes multidimensional generalisation of error bars
- ▶ Sampling is traditionally accomplished using random-walk MCMC tools like Gibbs sampling, Metropolis-Hastings, emcee or Hamiltonian Monte Carlo.



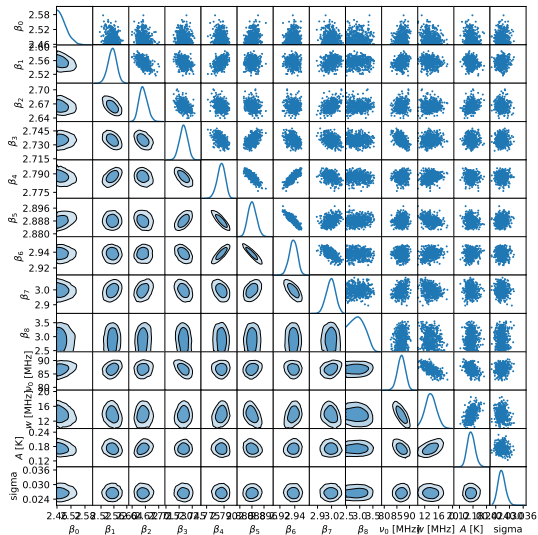
Numerical tools

- ▶ The key concept in numerical Bayesian inference is *sampling* a distribution
- ▶ i.e. drawing parameter vectors θ in proportion to the probability density $P(\theta)$.
- ▶ Compression of high-dimensional function.
- ▶ Posterior $P(\theta|D)$ encodes multidimensional generalisation of error bars
- ▶ Sampling is traditionally accomplished using random-walk MCMC tools like Gibbs sampling, Metropolis-Hastings, emcee or Hamiltonian Monte Carlo.




Numerical tools

- ▶ The key concept in numerical Bayesian inference is *sampling* a distribution
- ▶ i.e. drawing parameter vectors θ in proportion to the probability density $P(\theta)$.
- ▶ Compression of high-dimensional function.
- ▶ Posterior $P(\theta|D)$ encodes multidimensional generalisation of error bars
- ▶ Sampling is traditionally accomplished using random-walk MCMC tools like Gibbs sampling, Metropolis-Hastings, emcee or Hamiltonian Monte Carlo.




Nested sampling in action

- ▶ Nested sampling is a completely different way of sampling.
- ▶ Uses ensemble sampling to compress prior to posterior.
- ▶ Maintain a set S of n samples, which are sequentially updated:
 - S_0 : Generate n samples uniformly over the space (from the prior π).
 - S_{n+1} : Delete the lowest likelihood sample in S_n , and replace it with a new uniform sample with higher likelihood
- ▶ Requires one to be able to sample uniformly within a region, subject to a *hard likelihood constraint*.




Nested sampling in action

- ▶ Nested sampling is a completely different way of sampling.
- ▶ Uses ensemble sampling to compress prior to posterior.
- ▶ Maintain a set S of n samples, which are sequentially updated:
 - S_0 : Generate n samples uniformly over the space (from the prior π).
 - S_{n+1} : Delete the lowest likelihood sample in S_n , and replace it with a new uniform sample with higher likelihood
- ▶ Requires one to be able to sample uniformly within a region, subject to a *hard likelihood constraint*.




Nested sampling in action

- ▶ Nested sampling is a completely different way of sampling.
- ▶ Uses ensemble sampling to compress prior to posterior.
- ▶ Maintain a set S of n samples, which are sequentially updated:
 - S_0 : Generate n samples uniformly over the space (from the prior π).
 - S_{n+1} : Delete the lowest likelihood sample in S_n , and replace it with a new uniform sample with higher likelihood
- ▶ Requires one to be able to sample uniformly within a region, subject to a *hard likelihood constraint*.




Nested sampling in action

- ▶ Nested sampling is a completely different way of sampling.
- ▶ Uses ensemble sampling to compress prior to posterior.
- ▶ Maintain a set S of n samples, which are sequentially updated:
 - S_0 : Generate n samples uniformly over the space (from the prior π).
 - S_{n+1} : Delete the lowest likelihood sample in S_n , and replace it with a new uniform sample with higher likelihood
- ▶ Requires one to be able to sample uniformly within a region, subject to a *hard likelihood constraint*.




Nested sampling in action

- ▶ Nested sampling is a completely different way of sampling.
- ▶ Uses ensemble sampling to compress prior to posterior.
- ▶ Maintain a set S of n samples, which are sequentially updated:
 - S_0 : Generate n samples uniformly over the space (from the prior π).
 - S_{n+1} : Delete the lowest likelihood sample in S_n , and replace it with a new uniform sample with higher likelihood
- ▶ Requires one to be able to sample uniformly within a region, subject to a *hard likelihood constraint*.




Nested sampling in action

- ▶ Nested sampling is a completely different way of sampling.
- ▶ Uses ensemble sampling to compress prior to posterior.
- ▶ Maintain a set S of n samples, which are sequentially updated:
 - S_0 : Generate n samples uniformly over the space (from the prior π).
 - S_{n+1} : Delete the lowest likelihood sample in S_n , and replace it with a new uniform sample with higher likelihood
- ▶ Requires one to be able to sample uniformly within a region, subject to a *hard likelihood constraint*.




Nested sampling in action

- ▶ Nested sampling is a completely different way of sampling.
- ▶ Uses ensemble sampling to compress prior to posterior.
- ▶ Maintain a set S of n samples, which are sequentially updated:
 - S_0 : Generate n samples uniformly over the space (from the prior π).
 - S_{n+1} : Delete the lowest likelihood sample in S_n , and replace it with a new uniform sample with higher likelihood
- ▶ Requires one to be able to sample uniformly within a region, subject to a *hard likelihood constraint*.




Nested sampling in action

- ▶ Nested sampling is a completely different way of sampling.
- ▶ Uses ensemble sampling to compress prior to posterior.
- ▶ Maintain a set S of n samples, which are sequentially updated:
 - S_0 : Generate n samples uniformly over the space (from the prior π).
 - S_{n+1} : Delete the lowest likelihood sample in S_n , and replace it with a new uniform sample with higher likelihood
- ▶ Requires one to be able to sample uniformly within a region, subject to a *hard likelihood constraint*.




Nested sampling in action

- ▶ Nested sampling is a completely different way of sampling.
- ▶ Uses ensemble sampling to compress prior to posterior.
- ▶ Maintain a set S of n samples, which are sequentially updated:
 - S_0 : Generate n samples uniformly over the space (from the prior π).
 - S_{n+1} : Delete the lowest likelihood sample in S_n , and replace it with a new uniform sample with higher likelihood
- ▶ Requires one to be able to sample uniformly within a region, subject to a *hard likelihood constraint*.




Nested sampling in action

- ▶ At the end, one is left with a set of discarded points
- ▶ These may be weighted to form posterior samples
- ▶ They can also be used to calculate the evidence
- ▶ The evolving ensemble of live points allows algorithms to perform self-tuning and mode clustering




Nested sampling in action

- ▶ At the end, one is left with a set of discarded points
- ▶ These may be weighted to form posterior samples
- ▶ They can also be used to calculate the evidence
- ▶ The evolving ensemble of live points allows algorithms to perform self-tuning and mode clustering




Implementations of Nested Sampling


MultiNest




PolyChord



DNest



NeuralNest



Nested Sampling: Benefits and drawbacks

Relative to traditional numerical posterior samples (Metropolis Hastings, HMC, emcee), nested sampling:

- + Can calculate evidence (and therefore perform model comparison).
- + Can handle multi-modal distributions.
- + Requires little tuning for an a-priori unseen problem.
- + Highly parallelisable ($n_{\text{cores}} \sim n_{\text{live}} \gg 4$).
- Slower than a well-tuned posterior sampler.
- Run time is dependent on prior choice, and priors must be proper (some people view this as a feature rather than a bug).

Nested Sampling: a user's guide

1. Nested sampling is a likelihood scanner, rather than posterior explorer.
 - ▶ This means typically most of its time is spent on burn-in rather than posterior sampling
 - ▶ Changing the stopping criterion from 10^{-3} to 0.5 does little to speed up the run, but can make results very unreliable
2. The number of live points n_{live} is a resolution parameter.
 - ▶ Run time is linear in n_{live} , posterior and evidence accuracy goes as $\frac{1}{\sqrt{n_{\text{live}}}}$.
 - ▶ Set low for exploratory runs $\sim \mathcal{O}(10)$ and increased to $\sim \mathcal{O}(1000)$ for production standard.
3. Most algorithms come with additional reliability parameter(s).
 - ▶ e.g. MultiNest: eff , PolyChord: n_{repeats}
 - ▶ These are parameters which have no gain if set too conservatively, but increase the reliability
 - ▶ Check that results do not degrade if you reduce them from defaults, otherwise increase.

Key tools

- anesthetic** Nested sampling post processing [1905.04768]
github.com/williamjameshandley/anesthetic
- nestcheck** cross-checks using unthreaded runs [1804.06406]
github.com/ejhigson/nestcheck
- MultiNest** Ellipsoidal rejection sampling [0809.3437]
github.com/farhanferoz/MultiNest
- PolyChord** Python/C++/Fortran state of the art [1506.00171]
github.com/PolyChord/PolyChordLite
- dynesty** Python re-implementation of several codes [1904.02180]
github.com/joshspeagle/dynesty